

```

1 # imports
2 import pandas as pd
3 import numpy as np
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.model_selection import GridSearchCV
6 from sklearn.metrics import accuracy_score, f1_score, roc_auc_score,
  precision_score, recall_score
7 from bootstrap import bootstrap
8
9 # runs 5 trials of logistic regression on given dataset
10 def run_logistic_regression(X, y, n_trials=5):
11     """
12     Runs logistic regression for a dataset
13     5 trials of LogReg with 5-fold cross validation and a gridsearch over
  hyperparameter: C
14     and computes mean accuracy over metrics
15         accuracy, f1, roc, precision, recall
16
17     parameters
18     -----
19     X: feature vector
20     y: target vector
21     n_trials: number of trials to run
22
23     returns
24     -----
25     train_metrics: average of each metric on training set across 5 trials
26     test_metrics: average of each metric on test set across 5 trials
27     hyperp: dataframe of hyperparameters tried during hyperparameter search,
28             of metrics (acc, f1, roc) and their mean performance during
  cross-validation
29     """
30
31     # hyperparameters
32     C_list = [1e-8, 1e-7, 1e-6, 1e-5, 1e-4, 1e-3, 1e-2, 1e-1, 1e0, 1e1, 1e2, 1e3, 1e4]
33     params = {'C': C_list}
34
35     # metric evaluation scores
36     scores = ['accuracy', 'f1', 'roc_auc', 'precision', 'recall']
37
38     # to hold calculated metric performances
39     train_metrics = []
40     test_metrics = []
41
42     for trial in range(n_trials):
43
44         # initialize model for cross validation grid search
45         LogReg = LogisticRegression(max_iter=200)
46         GS = GridSearchCV(LogReg, params, scoring=scores, refit=False)
47

```

```

47
48     # bootstrap training and testing sets
49     X_train, X_test, y_train, y_test = bootstrap(X,y)
50     GS.fit(X_train, y_train)
51
52     # collect results and get dataframe to visualize hyperparameter
search
53     res = GS.cv_results_
54     hyperp = pd.DataFrame(res['params'])
55     hyperp['acc'] = res['mean_test_accuracy']
56     hyperp['f1'] = res['mean_test_f1']
57     hyperp['roc'] = res['mean_test_roc_auc']
58     hyperp.set_index('C', inplace=True)
59
60     test_per = [] # test set performances
61     train_per = [] # train set performances
62
63     # get best hyperparameters for each metric and use on test set
64     for s in scores:
65
66         # train logreg with best hyperparameters for metric
67         best_C = res['params'][np.argmax(res['mean_test_{}'.format(s)])]
68         LR = LogisticRegression(C=best_C['C'])
69         LR.fit(X_train, y_train)
70
71         # predictions for train and test sets
72         y_pred = LR.predict(X_test)
73         y_pred_train = LR.predict(X_train)
74
75         # evaluate metric on test set
76         if s == 'accuracy':
77             test_per.append(accuracy_score(y_test, y_pred))
78             train_per.append(accuracy_score(y_train, y_pred_train))
79         elif s == 'f1':
80             test_per.append(f1_score(y_test, y_pred))
81             train_per.append(f1_score(y_train, y_pred_train))
82         elif s == 'roc_auc':
83             test_per.append(roc_auc_score(y_test, y_pred))
84             train_per.append(roc_auc_score(y_train, y_pred_train))
85         elif s == 'precision':
86             test_per.append(precision_score(y_test, y_pred))
87             train_per.append(precision_score(y_train, y_pred_train))
88         elif s == 'recall':
89             test_per.append(recall_score(y_test, y_pred))
90             train_per.append(recall_score(y_train, y_pred_train))
91
92     train_metrics.append(train_per)
93     test_metrics.append(test_per)
94
95     print('Trial {} done'.format(trial+1))

```

```
96  
97     # take mean of each metric across 5 trials  
98     train_metrics = np.mean(np.array(train_metrics), axis=0)  
99     test_metrics = np.mean(np.array(test_metrics), axis=0)  
100  
101     return train_metrics, test_metrics, hyperp
```